

# Using several classifiers and tuning parameters - Parameters grid

[From official scikit-learn documentation](#)

Adapted by Claudio Sartori

Example of usage of the **model selection** features of `scikit-learn` and comparison of several classification methods.

1. import a sample dataset
2. do the usual preliminary data explorations and separate the predicting attributes from the *target* 'Exited'
3. define the *models* that will be tested and prepare the *hyperparameter ranges* for the modules
4. set the list of *score functions* to choose from
5. split the dataset into two parts: train and test
6. Loop on score functions and, for each score, loop on the model labels (see details below)
  - optimize with GridSearchCV
  - test
  - store the results
7. for each scoring show the ranking of the models, and the confusion matrix given by the best model
8. for each scoring show the confusion matrix of the prediction given by the best model

```
In [1]: """
@author: scikit-learn.org and Claudio Sartori
"""

import warnings
warnings.filterwarnings('ignore') # uncomment this line to suppress warnings

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn import datasets
from sklearn.model_selection import train_test_split, GridSearchCV, StratifiedKFold
from sklearn.metrics import classification_report, ConfusionMatrixDisplay
from sklearn.svm import SVC
from sklearn.linear_model import Perceptron
from sklearn.neural_network import MLPClassifier
```

```

from sklearn.tree import DecisionTreeClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import AdaBoostClassifier, RandomForestClassifier

print(__doc__) # print information included in the triple quotes at the beginning

```

@author: scikit-learn.org and Claudio Sartori

## 0. Initial settings

Set the random state and set the seed with `np.random.seed()`

Set the test set size and the number of cross validation splits

```

In [2]: ts = 0.3 # test size
        random_state = 42
        np.random.seed(random_state) # this sets the random sequence. Setting only this the repeatability is guaranteed
                                       # only if we re-execute the entire notebook

        random_state=None
        cv = 3 # number of cross-validation splits

```

## 1. Import the dataset

```

In [3]: url = 'churn-analysis.csv'
        df = pd.read_csv(url)
        df.head()

```

```

Out[3]:

```

|   | CreditScore | Gender | Age | Tenure | Balance   | NumOfProducts | HasCrCard | IsActiveMember | EstimatedSalary | Exited |
|---|-------------|--------|-----|--------|-----------|---------------|-----------|----------------|-----------------|--------|
| 0 | 619         | 0      | 42  | 2      | 0.00      | 1             | 1         | 1              | 101348.88       | True   |
| 1 | 502         | 0      | 42  | 8      | 159660.80 | 3             | 1         | 0              | 113931.57       | True   |
| 2 | 699         | 0      | 39  | 1      | 0.00      | 2             | 0         | 0              | 93826.63        | False  |
| 3 | 822         | 1      | 50  | 7      | 0.00      | 2             | 1         | 1              | 10062.80        | False  |

|   | CreditScore | Gender | Age | Tenure | Balance   | NumOfProducts | HasCrCard | IsActiveMember | EstimatedSalary | Exited |
|---|-------------|--------|-----|--------|-----------|---------------|-----------|----------------|-----------------|--------|
| 4 | 501         | 1      | 44  | 4      | 142051.07 | 2             | 0         | 1              | 74940.50        | False  |

## 2. Explore the data

The output of exploration is not shown here

```
In [4]: target = 'Exited'
X = df.drop(target, axis=1)#dataset.data
y = df[target]
```

## 3. Define the *models*

Prepare the *hyperparameter ranges* for the modules

Put everything in a dictionary, for ease of use

```
In [5]: model_lbls = ['dt' # decision tree
                    , 'nb' # gaussian naive bayes
                    # , 'lp' # linear perceptron
                    # , 'svc' # support vector
                    # , 'knn' # k nearest neighbours
                    # , 'adb' # adaboost
                    # , 'rf' # random forest
                    ]

models = {
    'dt': {'name': 'Decision Tree',
           'estimator': DecisionTreeClassifier(random_state=random_state),
           'param': [{'max_depth': [*range(1,20)], 'class_weight': [None, 'balanced']}]},
    'nb': {'name': 'Gaussian Naive Bayes',
           'estimator': GaussianNB(),
           'param': [{'var_smoothing': [10**exp for exp in range(-3,-12,-1)]]}
}
```

```

'lp': {'name': 'Linear Perceptron ',
      'estimator': Perceptron(random_state=random_state),
      'param': [{ 'early_stopping': [True, False], 'class_weight': [None, 'balanced'] }],
},
'svc': {'name': 'Support Vector ',
      'estimator': SVC(random_state=random_state),
      'param': [{ 'kernel': ['rbf'],
                    'gamma': [1e-3, 1e-4],
                    'C': [1, 10, 100],
                  },
                { 'kernel': ['linear'],
                    'C': [1, 10, 100],
                  },
                ],
},
},
'knn': {'name': 'K Nearest Neighbor ',
      'estimator': KNeighborsClassifier(),
      'param': [{ 'n_neighbors': list(range(1,7)) }],
},
'adb': {'name': 'AdaBoost ',
      'estimator': AdaBoostClassifier(random_state=random_state),
      'param': [{ 'n_estimators': [20, 30, 40, 50],
                    'learning_rate': [0.5, 0.75, 1, 1.25, 1.5] }],
},
'rf': {'name': 'Random forest ',
      'estimator': RandomForestClassifier(random_state=random_state),
      'param': [{ 'max_depth': [*range(4,10)],
                    'n_estimators': [*range(10,60,10)] }],
},
}
}

```

## 4. Set the list of *score functions* to choose from

```
In [6]: scorings = ['accuracy', 'precision_macro', 'recall_macro', 'f1_macro']
```

```
In [7]: # for development only
from sklearn.model_selection import ParameterGrid
trials = {lbl: len(list(ParameterGrid(models[lbl]['param']))) for lbl in models.keys()}
trials
```

```
Out[7]: {'dt': 38, 'nb': 9, 'lp': 4, 'svc': 9, 'knn': 6, 'adb': 20, 'rf': 30}
```

## 5. Split the dataset into the train and test parts

- the *\*train\** part will be used for training and cross-validation (i.e. for *\*development\**)
- the *\*test\** part will be used for test (i.e. for *\*evaluation\**)
- the fraction of test data will be `_ts_` (a value of your choice between 0.2 and 0.5)

```
In [8]: X_train, X_test, y_train, y_test = train_test_split(
        X, y, test_size=ts)
print("Training on {} examples".format(len(X_train)))
```

Training on 3509 examples

## 6. Loop on scores and, for each score, loop on the model labels

The function `GridSearchCV` iterates a cross validation experiment to train and test a model with different combinations of parameter values

- for each parameter we have set before a list of values to test, `ParametersGrid` will be implicitly called to generate all the combinations
- we choose a *score function* which will be used for the optimization
  - e.g. `accuracy_score`, `precision_score`, `recall_score`, `f1_score`, see this [page](#) for reference
- the output is a dataframe containing
  - the set of parameters maximising the score
  - the score used for optimisation and all the test scores

### Steps

- prepare an empty list `clfs` to store all the fitted models
- prepare an empty DataFrame which will collect the results of the fittings with each combination of parameters
  - dataframe columns are
 

```
'scoring', 'model', 'best_params', 'accuracy', 'precision_macro', 'recall_macro', 'f1_macro'
```
- loop

```
In [9]: clfs = []
results = pd.DataFrame(columns=['scoring', 'model', 'best_params', '#', 'fit+score_time',
                                'accuracy', 'precision_macro', 'recall_macro', 'f1_macro'])
```

## Parameters to collect

`classification_report` produces a dictionary containing some classification performance measures, given the *ground truth* and the *predictions* (use the parameter `output_dict=True` )

The measures are (among others):

- accuracy
- macro\_avg a dictionary containing:
  - precision
  - recall
  - f1-score
- ...

## Loop

- repeat for all the chosen scorings
  - repeat for all the chosen classification models
    - store in `clf` the initialisation of `GridSearchCV` with the appropriate
      - classification model
      - parameters ranges
      - scoring
      - cross validation method `cv` (the same for all)
    - fit `clf` with the *train* part of `X` and `y`
    - store in `y_pred` the prediction for the *test* part of `X`
    - append `clf` to `clfs`
    - append `y_pred` to `y_preds`
    - store in variable `cr` the `classification_report` produced with the test part of `y` and `y_pred`
    - store in the last row of `results` a list containing:
      - the name of the model

- the `.best_params_` of `clf`
- a selection of the contents of `cr`
  - 'accuracy',
  - 'macro avg'precision'
  - 'macro avg'recall'
  - 'macro avg'f1-score'

Hints:

- `cr` is a multi-level dictionary, second level can be reached with  
`cr['first level label']['second level label']`
- to append a list as the last row of a dataframe you can use  
`df.loc[len(df)]=[]`

```
In [11]: for scoring in scorings:
          for m in model_lbls:
              clf = GridSearchCV(models[m]['estimator'], models[m]['param'], cv=cv,
                                  scoring = scoring,

                                  )

              clf.fit(X_train, y_train)
              clfs.append(clf)
              y_true, y_pred = y_test, clf.predict(X_test)
              # y_preds.append(y_pred)
              cr = classification_report(y_true,y_pred, output_dict=True
                                         , zero_division=1
                                         )
              results.loc[len(results)] = [scoring,models[m]['name'],clf.best_params_
                                           # ,(clf.cv_results_['mean_fit_time'].sum()+clf.cv_results_['mean_score_time'].sum()
                                           ,cr['accuracy']
                                           ,cr['macro avg']['precision']
                                           ,cr['macro avg']['recall']
                                           ,cr['macro avg']['f1-score']]
```

## 7. Display

For each scoring show the ranking of the models, and the confusion matrix given by the best model

For each scoring:

- set a `scoring_filter`
- filter the results of that scoring
- display the filtered dataframe with the `display()` function (it allows several displays of dataframes )

In [12]:

```
for score in scorings:
    scoring_filter = score
    display(results[results.scoring==scoring_filter]\
            .sort_values(by=scoring_filter,ascending=False)\
            .drop('scoring',axis=1)\
            .style.format(precision=3)\
            .set_caption('Results for scoring "{}"'.format(scoring_filter))
    )
```

Results for scoring "accuracy"

|   | model                | best_params                            | accuracy | precision_macro | recall_macro | f1_macro |
|---|----------------------|--|----------|-----------------|--------------|----------|
| 0 | Decision Tree        | {'class_weight': None, 'max_depth': 4} | 0.872    | 0.800           | 0.700        | 0.734    |
| 1 | Gaussian Naive Bayes | {'var_smoothing': 1e-11}               | 0.847    | 0.811           | 0.574        | 0.588    |

Results for scoring "precision\_macro"

|   | model                | best_params                            | accuracy | precision_macro | recall_macro | f1_macro |
|---|----------------------|--|----------|-----------------|--------------|----------|
| 3 | Gaussian Naive Bayes | {'var_smoothing': 1e-11}               | 0.847    | 0.811           | 0.574        | 0.588    |
| 2 | Decision Tree        | {'class_weight': None, 'max_depth': 4} | 0.872    | 0.800           | 0.700        | 0.734    |

Results for scoring "recall\_macro"

|   | model                | best_params                                  | accuracy | precision_macro | recall_macro | f1_macro |
|---|----------------------|--|----------|-----------------|--------------|----------|
| 4 | Decision Tree        | {'class_weight': 'balanced', 'max_depth': 4} | 0.724    | 0.649           | 0.744        | 0.651    |
| 5 | Gaussian Naive Bayes | {'var_smoothing': 1e-11}                     | 0.847    | 0.811           | 0.574        | 0.588    |

Results for scoring "f1\_macro"

|   | model         | best_params                            | accuracy | precision_macro | recall_macro | f1_macro |
|---|---------------|--|----------|-----------------|--------------|----------|
| 6 | Decision Tree | {'class_weight': None, 'max_depth': 4} | 0.872    | 0.800           | 0.700        | 0.734    |



|   | model                | best_params              | accuracy | precision_macro | recall_macro | f1_macro |
|---|----------------------|--------------------------|----------|-----------------|--------------|----------|
| 7 | Gaussian Naive Bayes | {'var_smoothing': 1e-11} | 0.847    | 0.811           | 0.574        | 0.588    |

## 8. Confusion matrices

Use the `ConfusionMatrixDisplay` with the best model of each scoring to compare the predictions

Repeat for every scoring:

- filter the results for the current scoring
- find the row with the best value of the scoring; this row is also the index of the corresponding
- 

In [13]:

```
for score in scorings:
    scoring_filter = score
    # bests[score] = results.loc[results.scoring==scoring_filter,scoring_filter].idxmax(axis=0)
    best_row = results.loc[results.scoring==scoring_filter,scoring_filter].idxmax(axis=0)
    disp = ConfusionMatrixDisplay.from_estimator(X=X_test, y=y_test, estimator = clfs[best_row])
    # disp.ax_.set_title("Best Model for {}: {}".format(score,results.at[bests[score], 'model']))
    disp.ax_.set_title("Best Model for {}: {}".format(score,results.at[best_row, 'model']))
```







